

Thread Scheduling

Name - Kunal Sharma
Roll Number - 2021331
Branch - CSD

Code Description

I initially initialized the three thread IDs `pid1`, `pid2`, and `pid3` once starting with the main function. Then I used `pthread_create()` to create three new threads, each of which accepts a thread ID and a function (`CountA`, `CountB`, and `CountC`, respectively), each of which does the same counting operation from 1 to 2^{32} . I use the `clock_gettime()` function in these routines to determine how long it takes the function to complete its duty by using the attribute `CLOCK_REALTIME`.

I'm then using `pthread_setschedparam()`, which accepts the thread ID of the current thread and the policy (`SCHED_OTHER` for `CountA`, `SCHED_RR` for `CountB`, and `SCHED_FIFO` for `CountC`) as an attribute, to set the priority of the presently executing thread. Last but not least, I'm using `pthread_join()`, which waits for each thread to finish before executing context switching between various threads according to their priority, with the exception of the thread with policy `SCHED_OTHER`. I have successfully created 6 graphs by altering the priorities of the threads.

The output of the code shows a tendency for `SCHED_FIFO` to execute in the shortest amount of time and `SCHED_RR` to execute in a little bit more time because `SCHED_RR` performs poorly when the same task is being performed by these two threads (given that `SCHED_RR` and `SCHED_FIFO` have the same priority), while `SCHED_OTHER` executes much more slowly than `SCHED_RR` because it has the lowest priority by default and is executed last.