

Process Scheduling

Name - Kunal Sharma
Roll Number - 2021331
Branch - CSD

Code Description

In the main for loop, I first create three child processes that use the `exec()` command to run the three separate bash files (`bash1.sh`, `bash2.sh`, and `bash3.sh`). The actions that take place in each bash file are to copy the kernel, navigate to the copied kernel directory, and then compile the copied kernel using the `make` command. Using `sched_setscheduler()`, I set up a process policy, which is used as an attribute by the `sched_setscheduler`, and each child process uses `param.sched` priority to set the priority of the current process (`0`). Now I'm starting a for loop statement with the requirement that all of the child processes must be finished. I'm using `waitpid()` in this while loop to wait for the processes with their appropriate process IDs to complete. Finally, I'm including a while loop to figure out the time that `clock_gettime()` marked using the attribute `CLOCK_REALTIME`.

The output of the code shows a tendency for `SCHED_RR` to execute in the shortest amount of time and `SCHED_FIFO` to execute in a slightly longer amount of time because `SCHED_FIFO` performs poorly when the same task is being performed by these two processes (given priority of `SCHED_RR` and `SCHED_FIFO` is same) . `SCHED_OTHER` takes a lot longer than `SCHED_FIFO` because `SCHED_OTHER` has the lowest priority by default and is executed last.